

LOGESHWARAN SHANMUGAM

AI Engineer | LLM Systems Developer

@logeshwaranshanmugam02@gmail.com

+91 87788 20191

Bengaluru, India

logeshwaran-shanmugam

PROFILE

AI Engineer specializing in Large Language Models (LLMs), Agentic AI Systems, Retrieval-Augmented Generation (RAG), and AI workflow automation. Experienced in building scalable AI platforms, MCP servers, AI API gateways, and multi-agent systems using FastAPI, LangChain, LangGraph, and vector databases. Skilled in LLM orchestration, semantic search, tool calling, context engineering, and AI infrastructure development.

EXPERIENCE

AI/ML Engineer

IAgami Technologies

Jul 2025 – Present

- Designed and developed AI-powered backend services using FastAPI, LangChain, LangGraph, and LLM orchestration frameworks, cutting integration time for new AI features across teams.
- Developed MCP servers enabling standardized, tool-calling interaction between AI agents and enterprise systems.
- Built Agentic AI workflows integrating external APIs, enterprise tools, vector databases, and custom business logic.
- Implemented Retrieval-Augmented Generation (RAG) pipelines using ChromaDB, FAISS, and Neo4j, improving enterprise knowledge retrieval accuracy and response relevance.
- Engineered core routing and monitoring logic for AgamiLLM, an AI Gateway unifying **20+ LLM models** and handling **10K+ requests/month** with **99.5% availability** and **<100ms average latency**.
- Reduced provider integration effort by **~60%** across **3 internal teams** and **10+ active users** through centralized model management and a single unified API layer.
- Containerized AI services using Docker, improving deployment consistency and reducing environment-related issues across releases.
- Integrated Google Cloud Platform (GCP) services and external AI APIs into enterprise AI workflows.
- Built **15+ n8n automation workflows**, saving **10–15 hours of manual effort weekly** and cutting repetitive operational tasks by **~60%**.

Full Stack Developer Intern

IAgami Technologies

2024 – 2025

- Built full-stack applications using React.js and backend APIs, delivering features from design through deployment.
- Developed LLM-powered chatbot systems using retrieval-based knowledge pipelines to automate internal query handling.
- Implemented backend APIs supporting AI-driven applications, laying groundwork later extended into production RAG systems.

CORE SKILLS

AI Systems

Agentic AI, RAG, Multi-Agent Systems, MCP, Tool Calling, Function Calling, Prompt Engineering, Context Engineering, Semantic Search

Vector Databases

ChromaDB, FAISS, Neo4j

Backend & Databases

FastAPI, Spring Boot, REST APIs, PostgreSQL, MySQL

Cloud & DevOps

Google Cloud Platform (GCP), Docker

LLM Frameworks

LangChain, LangGraph, LiteLLM, OpenRouter

Programming & AI Libraries

Python, Java, JavaScript, SQL, Transformers, Scikit-learn, Pandas, SentenceTransformers, NumPy

Automation & Developer Tools

n8n, GitHub, VS Code, Postman, Jupyter Notebook

EDUCATION

B.Tech in Artificial Intelligence & Data Science

Sri Shanmugha College of Engineering & Technology

📅 2021 – 2025

- GPA: 8.2

PROJECTS

AI API Gateway

AgamiLLM

📅 2025 – Present

- Engineered an AI Gateway handling **10K+ LLM requests/month**, reducing provider integration effort by **~60%** while maintaining **99.5% availability** and **<100ms average gateway latency**.
- Integrated and managed **20+ AI models** through a unified API layer, giving developers a single interface instead of per-provider integrations.
- Supported AI access for **3 teams and 10+ internal users** through centralized model management.
- Implemented model routing, authentication, usage tracking, monitoring, and provider abstraction.
- Containerized gateway services using Docker for deployment and scalability.

RAG-based Support Assistant

IAgami Platform

📅 Nov 2025 – Dec 2025

- Developed a RAG-based Support Assistant that automated **1,500+ queries/month**, resolving **~65%** of common queries without human escalation and reducing resolution time by **~40%**.
- Built document ingestion pipelines, embedding generation, and vector search.
- Integrated ChromaDB, FAISS, and Neo4j for enterprise knowledge retrieval.
- Enabled context-aware responses through prompt context injection.

HR Knowledge Assistant

Enterprise HR System

📅 2025

- Built an HR Knowledge Assistant that automated **500+ queries/month**, reducing repetitive HR ticket volume by **~30%** and delivering responses within seconds.
- Implemented tool-calling workflows allowing employees to apply for leave, view leave balances, and track leave requests directly from chat.
- Built LangGraph reasoning workflows to evaluate user intent and dynamically execute appropriate tools.
- Integrated leave application, leave status retrieval, and HR policy assistance into a unified conversational experience.
- Developed FastAPI services and enterprise integrations supporting real-time HR operations.

NutriLLM – AI Nutritionist Platform

Agentic AI System

📅 Mar 2026 – Jun 2026

- Developed NutriLLM to generate **1,000+ personalized diet plans**, reducing plan-generation time from 15–20 minutes manually to **under 1 minute** and cutting manual effort by **~90%**.
- Designed a meal planning engine that dynamically adjusts food quantities to match target nutritional ranges.
- Integrated nutrition datasets and food databases for calorie and macronutrient calculations.
- Implemented AI-driven reasoning workflows for personalized meal recommendations.

ACHIEVEMENTS

- Received the **Top Academic Performer Award** from the Department of AI & DS at Sri Shanmugha College.
- Coordinated and organized a successful technical symposium event.
- Currently upskilling in advanced agentic workflows via LangChain Academy's *Intro to LangGraph* course.